

Computer Tutorial 1: Matlab Basics

Exercise 1: *Basic Matrix Commands*

A sample program

%This is a simple program which illustrates Matlab matrix commands

```
x = [1 2 3; 4 5 6; 7 8 9];
```

```
disp x;
```

```
disp(x);
```

```
y = [11 12 13; 14 15 16; 17 18 19];
```

```
disp y;
```

```
disp(y);
```

```
z=x+y;
```

```
disp z ;
```

```
disp(z);
```

```
w=x-y;
```

```
disp w;
```

```
disp(w);
```

```
u= x*y;
```

```
disp u;
```

```
disp(u);
```

```
a=[x, y];
```

```
disp a;
```

```
disp(a);
```

```
b=[x; y];
```

```
disp b;
```

```
disp(b);
```

```
c=x(:,2);
```

```
disp c;
```

```
disp(c);
```

```
d=y(:,1);
```

```
disp d;
```

```
disp(d);
```

```
e=x(2,3);
```

```
disp e
```

```
disp(e);
```

- Create the preceding program and run it in Matlab and examine the output Matlab produces. Describe what each line of this program does.
- Add a line to this program which creates a new matrix, f, which is the transpose of x.
- Add a line to this program which creates a new matrix, g, which is the identity matrix plus x.
- Add a line to this program which creates a new matrix, ginv, which is the inverse of g. What happens if you try to take the inverse of x?

Exercise 2: *OLS Estimation Using Artificial Data*

Matlab has many scripts (i.e. built-in little programs) that you can call automatically as part of your program. (Note: You can also create your own scripts). Here we show you how to use Matlab's scripts for random number generation from the Normal and Uniform distributions to create an artificial data set from the regression model: $y_i = \alpha + \beta x_i + \varepsilon_i$ for $i = 1, \dots, 100$. We set $\alpha = 1$, $\beta = 2$ and let the x_i and ε_i be random draws from the $U(0, 1)$ and $N(0, 1)$ distribution, respectively.

A sample program

%This is a program which artificially creates a data set and then does OLS estimation using it

%First part of this program artificially simulates data

```
n=100;
```

```
alpha=1;
```

```
beta=2;
```

```

e = randn(n,1);
x=rand(n,1);
y=alpha + x*beta + e;
%following line adds intercept to x. explain why
x=[ones(n,1), x];
%Following part of the program does OLS estimation
bhat = inv(x'*x)*x'*y;
disp 'The OLS estimate of beta is';
disp(bhat);
%the OLS residuals
resids = y - x*bhat;
%The OLS estimate of the error variance
s2 = resids'*resids/(n-2);
disp 'The OLS estimator of the error variance is';
disp(s2);

```

- Create this program and run it in Matlab and examine the output in Matlab. Describe what each line of this program does.
- Extend this program to calculate the R^2 of this regression and print out the result.
- Extend this program to calculate the covariance matrix of the OLS estimators (i.e. $\text{var}(\hat{\beta}_{OLS}) = s^2 (X'X)^{-1}$) and print out the result.

Exercise 3: For Loops and If Statements

When doing Monte Carlo integration or Gibbs sampling we repeatedly take draws from the posterior distribution. Matlab does this kind of repeated action using constructs called "for loops".

A sample program

```

%This is a program which illustrates for loops and if statements
%first create a column vector to work with
x=[1;2;7;5;9;3;6;9;1;11;1];
%the following command sums up the elements of a column vector
xsum=sum(x);
xsum1=0;
for i=1:11
xsum1=xsum1 + x(i,1);
end
disp xsum;
disp(xsum);
disp xsum1;
disp(xsum1);
%now illustrate the if command
xsum2=0;
for i=1:11
if x(i,1)>4
xsum2=xsum2 + x(i,1);
end
end
disp xsum2
disp(xsum2);

```

- Create this program and run it in Matlab and examine the output. Describe what each line of this program does. In particular, why are xsum1 and xsum the same as one another? What does the "if" statement do? What is "xsum2"?
- The sample program *sums* various column vectors. Modify this program to calculate *averages* (i.e. means).

Exercise 4: Drawing from Standard Distributions

Simulation-based inference using algorithms such as the Gibbs sampler requires the researcher to be able to draw from standard distributions. In this exercise we discuss how MATLAB can be used to obtain draws from a variety of standard continuous distributions. Specifically, we obtain draws from the Uniform, Normal, Student-t, Beta, Exponential and Chi-squared distributions (see the Appendix of Bayesian Econometric Methods for definitions of these distributions). Using the Matlab program for this exercise (Ex4.m), obtain sets of 10, 100 and 100,000 draws from the Uniform, standard Normal, Student-t(3) (denoted $t(0, 1, 3)$ in the notation of the Appendix to the book), Beta(3,2), Exponential with mean 5 and $\chi^2(3)$ distributions. For each sample size calculate the mean and standard deviation and compare these quantities to the known means and standard deviations from each distribution.

Exercise 5: Monte Carlo Integration

If the posterior density $p(\theta|y)$ takes the a familiar form (e.g. a Normal or Student-t or Gamma or other distribution for which computer algorithms exist to take random draws) then we can obtain R i.i.d. draws of the parameters, which we denote $\theta^{(r)}$, $r = 1, \dots, R$. Usually, quantities of interest to the researcher are functions of the model parameters. Let us call such a function $g(\theta)$. The researcher would then often be interested in calculating:

$$E(g(\theta)|y) = \int g(\theta) p(\theta|y) d\theta$$

Monte Carlo integration allows us to calculate integrals of this form. The weak law of large numbers implies that

$$E(g(\theta)|y) \simeq \frac{\sum_{r=1}^R g(\theta^{(r)})}{R}$$

This means that the posterior mean of $g(\theta)$ can be calculated by drawing from the posterior and then averaging functions of the posterior draws. Exercise: Suppose $p(\theta|y) \sim N(1, 4)$ and the quantity of interest is $g(\theta) = \theta^2$. Use Monte Carlo integration to calculate $E(\theta^2)$. Code for this question is in Ex5.m. Note: in this case, you know the correct answer is $E(\theta^2) = 5$ (since the definition of variance tells you that $var(\theta) = E(\theta^2) - [E(\theta)]^2$ and, in this exercise, $var(\theta) = 4$ and $E(\theta) = 1$), so you would not need to have done Monte Carlo integration. Optional exercise: modify Ex5.m to calculate the posterior mean for a more complicated quantity of interest for which analytical results are not so easily available (e.g. calculate $\Pr(\theta^2 > 2)$ or $E(\ln(\theta))$ or some other choice for $g(\theta)$).